

Indexing and Query optimization


Hogeschool Rotterdam
Rotterdam, Netherlands

Lecture topics

- Query optimization.
- Examples of slow query operations.
- Hashing.
- Trees.

Reasons

- Query needs to be fast.
- Sometimes they are not.
- You do not want to see your nephew born before retrieving the book you are looking for from Amazon.



Causes

- Too much data (Big data analysis).
 - Data clustering.
 - Better hardware. (arrays of disks, caching, ...)
- Too complex queries (DBMS optimization)
 - Refactor query. (Access planner)
 - Refactor data. (Indexing)

Indexes

- Query refactoring not always possible.
- Build additional data to speed up the data retrieval.

Indexes

- Take your text book and look for the paragraph titled “Key constraints” without using the index. How many pages have you looked?

Indexes

- Take your text book and look for the paragraph titled “Key constraints” without using the index. How many pages have you looked?
- **Answer:** 29 (from page 3 to 32).

Indexes

- Take your text book and look for the paragraph titled “Key constraints” without using the index. How many pages have you looked?
- **Answer:** 29 (from page 3 to 32).
- Do the same using the index. How many pages have you looked?

Indexes

- Take your text book and look for the paragraph titled “Key constraints” without using the index. How many pages have you looked?
- **Answer:** 29 (from page 3 to 32).
- Do the same using the index. How many pages have you looked? **Answer:** 2 (1 in the index, 1 in the text).

```
SELECT name
FROM ships
WHERE firepower = 1500
```

ships				
name	type	firepower	speed	position
Red 1	X-Wing	10	300	(1,3,1)
Red 2	X-Wing	10	300	(1,2,1)
Red 3	X-Wing	10	300	(0,2,5,1)
Red 4	X-Wing	10	300	(2,2,5,1)
Red 5	X-Wing	10	300	(2,2,5,0)
Red 6	X-Wing	10	300	(1,2,5,0)
Tantine IV	Corellian Corvette	60	300	(4,2,5,0)
Tyranny	Imperial Star Destroyer	1500	100	(12,0,0)
Accuser	Imperial Star Destroyer	1500	100	(-12,0,0)
Bombard	Victory Star Destroyer	1500	175	(-6,1,0)

```
SELECT name
FROM ships
WHERE firepower = 1500
```

ships				
name	type	firepower	speed	position
Red 1	X-Wing	10	300	(1,3,1)
Red 2	X-Wing	10	300	(1,2,1)
Red 3	X-Wing	10	300	(0,2,5,1)
Red 4	X-Wing	10	300	(2,2,5,1)
Red 5	X-Wing	10	300	(2,2,5,0)
Red 6	X-Wing	10	300	(1,2,5,0)
Tantine IV	Corellian Corvette	60	300	(4,2,5,0)
Tyranny	Imperial Star Destroyer	1500	100	(12,0,0)
Accuser	Imperial Star Destroyer	1500	100	(-12,0,0)
Bombard	Victory Star Destroyer	1500	175	(-6,1,0)

Number of comparisons: 10

Indexes

- How many comparisons we do at most in a table with R records?

Indexes

- How many comparisons we do at most in a table with R records?
- **R comparisons**

Indexes

- How many comparisons we do at most in a table with R records?
- **R comparisons**
- How many comparisons we do at least in a table with R records?

Indexes

- How many comparisons we do at most in a table with R records?
- **R comparisons**
- How many comparisons we do at least in a table with R records?
- **R comparisons**

Indexes

- How many comparisons we do at most in a table with R records?
- **R comparisons**
- How many comparisons we do at least in a table with R records?
- **R comparisons**
- Selection always requires to scan the entire table.

- Sorting and grouping requires to sort the column values.
- The best sorting algorithm requires about $R \log R$ operations, where R is the number of records.
- Running the query below requires about $10 * \log 10 \simeq 23$ operations.

```
SELECT type
FROM ships
WHERE firepower = 500
ORDER BY name DESC
```

- Generate pairs with one element from the first table and the second from the other followed by a selection.
- Same problem of the selection.
- Consider the following query applied to ship and the table below.

```
SELECT s.name,p.damage  
FROM ships s,projectiles p  
WHERE s.position = p.position AND  
       s.name = p.target
```

Projectiles		
target	position	damage
Red 3	(0,1,0)	30
Red 3	(3,1,-2)	50
Red 3	(0,2.5,1)	100

JOIN performance

- How many comparisons does the join make?

JOIN performance

- **How many comparisons does the join make?**
- Each entity of the first table must be compared.
- For each entity of the first table there is a comparison with each entity of the second for each selection condition.
- **Total comparisons:** $2 \cdot 10 \cdot 3 = 60$.
- **How many operations does JOINING two tables with one condition, respectively with N and M records, require?**

JOIN performance

- **How many comparisons does the join make?**
- Each entity of the first table must be compared.
- For each entity of the first table there is a comparison with each entity of the second for each selection condition.
- **Total comparisons:** $2 \cdot 10 \cdot 3 = 60$.
- **How many operations does JOINING two tables with one condition, respectively with N and M records, require?**
- $N \cdot M$.
- **What if there are C conditions?**

JOIN performance

- **How many comparisons does the join make?**
- Each entity of the first table must be compared.
- For each entity of the first table there is a comparison with each entity of the second for each selection condition.
- **Total comparisons:** $2 \cdot 10 \cdot 3 = 60$.
- **How many operations does JOINING two tables with one condition, respectively with N and M records, require?**
- $N \cdot M$.
- **What if there are C conditions?**
- $C \cdot N \cdot M$.

Dictionaries and Indices

- A *Dictionary* is a collection of *entries*.
- An *entry* is a pair $\langle key, value \rangle$.
- An index can be thought as a dictionary. The key is the search parameter in the index.

Hash table

- It is an array of *buckets*. Each element of the array is a pointer to a specific bucket.
- A *bucket* is an array used to contain data.
- Insertion uses a *hash function* to map a key to a position in the array.
- The *hash function* always returns a value between 0 and `array.length - 1`.
- Use the hash function to search for a key and go to the corresponding bucket.

Hash table

Indexing and
Query
optimization

Introduction

Query
optimization

Problematic
queries

WHERE

SORTING

JOIN

Indexing

Hashing

Trees




Figure: Example of data insertion in a hash table

Collisions

- The hash function might return duplicate values for two different keys, or for duplicate keys.
- We add an element to the bucket for each collision.
- **What happens if the bucket is full? (*overflow*)**

Collisions

- The hash function might return duplicate values for two different keys, or for duplicate keys.
- We add an element to the bucket for each collision.
- **What happens if the bucket is full?** (*overflow*)
- Add a list as overflow bucket. All overflowing entries are stored in the overflow bucket.
- Extend the bucket array when needed.

Performance (WHERE)

- Consider the query on ships presented in the WHERE slides. Suppose we have a hash table on firepower.
- **How many operations does the selection require?**

Performance (WHERE)

- Consider the query on ships presented in the WHERE slides. Suppose we have a hash table on firepower.
- **How many operations does the selection require?**
- 1 to access the correct bucket, plus 3 operations to read the records in the same bucket = 4 operations vs 10 of non-indexed implementation.

Indexing and
Query
optimization

Introduction

Query
optimization

Problematic
queries

WHERE

SORTING

JOIN

Indexing

Hashing

Trees

- **How can the join be implemented?**

- **How can the join be implemented?**

```
for (p : projectiles)
  for (s : ships)
    if (s.name = p.target &&
        s.position = p.position)
      add <s.name,p.damage> to result
```

- We can index one of the tables and put the indexed query as inner query in the for-loop.

Performance (WHERE)

- What is the cost of the join with a hash index on `namea`?

Performance (WHERE)

- **What is the cost of the join with a hash index on `name`^a?**
- For each record in `projectile` we need to access the bucket in the other table
- The only record which matches in `ship` is 'Red 3'.
- We apply the hash function and read the entry in the hash table for 'Red 3'. This means $3 \cdot 2 = 6$ operations to select the ships in both table with the same name.
- In total we have $6 + 30$ operations = 36 operations (30 for the non-indexed attribute in the condition).

^aAssume that the index maps the starting letter of its position in the alphabet order and no collisions

Performance (WHERE)

- What if we also have an index on position?

Performance (WHERE)

- **What if we also have an index on position?**
- There is one entity in the `ships` table which matches the condition.
- For each record in `projectiles` we look for a record in `ships` with the same position.
- 2 records mismatch the position and one matches.
- Accessing the mismatching records requires 2 operations (we only need to hash the search key to find it is not in the buckets). Accessing the matching record requires 2 operations. In total 4 operations.
- With two indices we require only $6 + 4 = 10$ operations.

Drawbacks of hash tables

- Useless if the condition is not an equality but an inequality (conditions on intervals).
- Useless with ordering/grouping commands.
- A lot of collisions require to extend the buckets/slow down the search if overflow buckets are used.

Balanced trees

- A graph is a set of vertices connected by edges.
- A tree is a graph without cycles and connected (all nodes can be reached following connections from a starting node).
- The top node is called root.
- The nodes that are directly connected to a node and at a deeper level are called children. The children shares the same parent.
- A node which does not have children is called *leaf*.
- A balanced tree is a tree where the leaves are all at the same level.

2-3 Trees

- Balanced tree.
- Every node contains at most 2 entries
- Every node has 2 or 3 children.
- The keys in the left child are less than or equal to the keys in the parent.
- The keys in the middle child are between the min and the max keys stored in the parent.
- The keys in the right child are greater than or equal to the keys in the parent.

2-3 Trees

Indexing and
Query
optimization

Introduction

Query
optimization

Problematic
queries

WHERE

SORTING

JOIN

Indexing

Hashing

Trees




Figure: Index on firepower

Performance (WHERE)

- Consider the query on ships presented in the WHERE slides. Suppose we have a 2-3 tree index on firepower.
- **How many operations does the selection require?**

Performance (WHERE)

- Consider the query on ships presented in the WHERE slides. Suppose we have a 2-3 tree index on firepower.
- **How many operations does the selection require?**
- We follow the index and we find 3 entries, reading 4 nodes. We need 4 operations vs 10 operations without an index.

Performance (JOIN)

- **What is the cost of the join with a 2-3 tree index on name in table ships (see the next slide for the index) ^a?**

Performance (JOIN)

- **What is the cost of the join with a 2-3 tree index on name in table ships (see the next slide for the index) ^a?**
- This time the inner table in the for-loop is ships.
- We need to find 'Red 3' three times in the index. This requires to traverse the whole tree 3 times and each time we access the record, for a total of 12 operations.
- We have 12 operations vs the required 30 without indexing.
- In total we have 42 operations vs 60 (position is not indexed).
- This is a very unlucky case, since the entry is stored in a leaf. Imagine we were looking for "Red 4", we would need just 1 operation to get to the correct node.

^aAssume that the index maps the starting letter of its position in the

Performance (JOIN)

Indexing and
Query
optimization

Introduction

Query
optimization

Problematic
queries

WHERE

SORTING

JOIN

Indexing

Hashing

Trees




Figure: Index on **name**

Performance (JOIN)

- Can you guess what is the generalized formula for the number of operations to search for a key in a 2-3 tree?

Performance (JOIN)

- Can you guess what is the generalized formula for the number of operations to search for a key in a 2-3 tree?
- $\log_2 N$ where N is the number of entries.

- **What if we also have an index on position** (see the next slide for the index)¹?

¹Imagine that the key is generated by summing the components of the position

- **What if we also have an index on position** (see the next slide for the index)¹?
- The position will not match for 2 out of 3 records in projectile. So we have to traverse 3 nodes in the tree twice for a total of 6 operations.
- The last position will match and the entry is stored in the root so we need just 1 operation.
- In total we have 7 operations.
- With both indices we need 19 vs 60 operations without indices.

¹Imagine that the key is generated by summing the components of the position

Performance (JOIN)

Indexing and
Query
optimization

Introduction

Query
optimization

Problematic
queries

WHERE

SORTING

JOIN

Indexing

Hashing

Trees




Figure: Index on position

Ordering/grouping with trees

- A clustered index is an index where the record order in a file is the same in the index.
- If the tree index is clustered we can use the tree for ordering/grouping.
- Just traverse the tree. If the ordering is ascending then visit all the nodes starting from the leftmost and then moving to the rightmost.
- This requires N steps where N is the number of entries in the index (vs $N \log N$ of a traditional ordering algorithm).

Exercise: order the `name` attributes in ascending order.

Exercise: order the `name` attributes in ascending order.

Overall evaluation of trees

- Fast with conditions on intervals.
- Fast with ordering/group by.
- The worst case requires to scan $\log N$ entries before finding our entries (with Hash tables it is at most the size of the largest bucket).

Golden rules of indexing

- If most of the queries have selections/joins with equalities choose hash tables.
- If most of the queries have selections/joins with inequalities choose trees.
- If you can anticipate there will be a lot of mismatches in the comparisons choose hash tables.
- If you anticipate there will be a lot of collisions (a lot of duplicate values) choose trees.
- If you have a lot of queries with ordering/group by use clustered trees (otherwise do not use clustering).