



HOGESCHOOL ROTTERDAM / CMI

Advanced Databases (Development 6B)

INFDEV026B

Number of study points: 4 ects

Course owners: Ahmad Omar, Francesco Di Giacomo



Modulebeschrijving

Module name: (Development 6B)	Advanced Databases
Module code:	INFDEV026B
Study points and hours of effort for full-time students:	<p>This module gives 4 ects, in correspondance with 112 hours:</p> <ul style="list-style-type: none"> • 3 x 7 hours frontal lecture. • the rest is self-study for the theory and practicum.
Examination:	Written examination and practicum (with oral check)
Course structure:	Lectures, self-study, and practicum
Prerequisite knowledge:	None.
Learning tools:	<ul style="list-style-type: none"> • Book: Database management systems (3rd edition); authors Ramakrishnan and Gehrke • Book: NO SQL Distilled; authors Sadalage and Fowler • Presentations (in pdf): found on N@tschool and on the GitHub repository https://github.com/hogeschool/INFDEV03-5 • Assignments, to be done at home (pdf): found on N@tschool and on the GitHub repository https://github.com/hogeschool/INFDEV03-5
Connected to competences:	<ul style="list-style-type: none"> • Analysis, design, and realisation of software at level 2
Learning objectives:	<p>At the end of the course, the student can:</p> <ul style="list-style-type: none"> • realise a normalized relational database and implement an application to execute operations on it RDBMS, NORM • describe the differences between relational and non-relational databases NONREL • use the map-reduce paradigm to run queries. NONREL • describe models of concurrency and transactions in a modern DBMS TRANS



Content:	<ul style="list-style-type: none">• relevant concepts in relational databases normalization• fundamental properties of DBMS's: atomicity, consistency, isolation, and durability (ACID)• concurrency and transactions• Map-reduce paradigm.• BASE vs ACID.• Graph databases.
Course owners:	Ahmad Omar, Francesco Di Giacomo
Date:	12 november 2017



1 General description

Databases are ubiquitous within the field of ICT. Many business needs are centered around the gathering, elaboration, etc. of large amounts of data. This data is crucially connected to real-world operations and thus its constant availability and timely elaboration is of unmissable importance.

This course covers advanced aspects of data processing and elaboration within the different sets of tradeoffs of the precise but limiting ACID framework and the imprecise but forgiving BASE framework.

1.1 Relationship with other teaching units

This course builds upon the basic databases course.

Knowledge acquired through the databases course is also useful for some of the projects. A word of warning though: projects and development courses are largely independent, so some things that a student learns during the development courses are not used in the projects, some things that a student learns during the development courses are indeed used in the projects, but some things done in the projects are learned within the context of the project and not within the development courses.



2 Course program

The course is structured into six lecture units. The lecture units are not necessarily in a one-to-one correspondence with the course weeks.

2.1 Unit 1 - Review

The course starts with a quick review on SQL and RDBMS's:

Topics

- Entities and relationships
- SQL operators

2.2 Unit 2 - Normalization

Theoretical concepts behind normalization of relational models. Normalization algorithms:

Topics

- Redundancy problem.
- Definition of functional dependencies.
- Normal forms definition.

2.3 Unit 3 - Normalization algorithms

In this unit we cover the main normalization techniques.

Topics

- Normalization in 1NF, 2NF, 3NF, BCNF.
- Exercises on normalization.

2.4 Unit 4 - Concurrency

The fourth lecture covers handling of potentially conflicting concurrent query execution in an ACID DBMS:

Topics

- ACID property.
- Serialization
- Locks
- Deadlocks and their prevention

2.5 Unit 5 - NOSQL: Map-Reduce paradigm

The fifth lecture covers map-reduce paradigm:

Topics

- Map function.
- Reduce function.
- Map-Reduce is SELECT-FROM-WHERE
- Idea behind NoSQL
- LINQ and map-reduce in LINQ.



2.6 Unit 6 - Graph databases

The sixth lecture covers a specific example of no-SQL databases, specifically graph databases:

Topics

- Directed vs undirected graphs
- Adjacency list vs matrix
- Algorithms on graphs
- Case study: Neo4J



3 Assessment

The course is tested with two exams: a series of practical assignments, a brief oral check of the practical assignments, and a written exam. The final grade is determined as follows:

$0.6 * \text{practicum} + 0.4 * \text{written exam}$

To pass the exam you must have a positive (i.e. ≥ 5.5) grade in both parts.

Motivation for grade A professional software developer is required to be able to program code which is, at the very least, *correct*.

In order to produce correct code, we expect students to show: *i*) a foundation of knowledge about how a programming language actually works in connection with a simplified concrete model of a computer; *ii*) fluency when actually writing the code.

The quality of the programmer is ultimately determined by his actual code-writing skills. The quick oral check ensures that each student is able to show that his work is his own and that he has adequate understanding of its mechanisms. The theoretical exam tests that the required foundation of knowledge about databases is also present.

3.1 Theoretical examination

The general shape of a theoretical exam for the course is made up of a series of highly structured open questions. In each exam the content of the questions will change, but the structure of the questions will remain the same. For the structure (and an example) of the theoretical exam, see the appendix.

3.2 Practical examination

The assignments must be delivered within Sunday 23 October at 23:59. If the student fails to respect this deadline he will have to do the oral check during the retake. The assignments must be delivered on Natschool for all the students (even in the same group) in a .zip folder. Each assignment will have a separate delivery. The content of the delivery will be specified in the assignment description published during the course.

The score for each assignment is given in the following way:

- **Assignment 1:** 4 pts.
- **Assignment 2:** 3 pts.
- **Assignment 3:** 3 pts.



Bijlage 1: Toetsmatrijs

Learning goals	Dublin descriptors
RDBMS	1, 2, 3, 4, 5
NORM	1, 2, 5
TRANS	1, 4
NONREL	1, 2, 5

Dublin-descriptors:

1. Knowledge and understanding
2. Applying knowledge and understanding
3. Making judgements
4. Communication
5. Learning skills



4 Practicum assignments

In this section the practicum assignments and their grading criteria are listed.

Assignment 1 The student is asked to implement a simple database for a given real-world scenario. This database is managed by a management application written in a chosen OO-language (Java or C#). The student is also asked to provide analysis on possible functional dependencies and provide a normalized version of the database.

Assignment 2 Implement the scenario of the first assignment in a document-based database. The implementation must come with a management application for the database using the map-reduce paradigm to execute given queries and to populate the database.

Assignment 3 Given the scenario described in the previous assignments, the student is asked to realise a correspondent graph database model, and implement it in graph DBMS Neo4j. The database must be populated by using the Cypher language coming with Neo4j, as well as the queries.

Assignment grading:

The grading system is showed in the table below. The teacher reserves the right to assign half of the points if the goal is partially achieved. If the student is individually unable to answer a question about any of the assignments, the student will not receive the score for that assignment.

Note that in Assignment 1 you can receive an extra point for the ERD but the total score cannot exceed the maximum amount of 4 points. Also note that the front end application connected to a database is a minimum requirement for the assignment to be graded. If that is missing, the grade will be automatically 0.

Assignment 1	<ul style="list-style-type: none"> • All the data is correctly modelled by the database (1 pt.). • The report contains significant functional dependencies and an explanation about the normal form of each table, possibly showing normalization steps (1 pt.). • All the operations listed in the assignment description are correctly implemented (1 pt.). • The application uses design patterns to abstract the database layer from the application layer (1 pt.) • There is a correct ERD modelling the database (extra 1 pt.)
Assignment 2	<ul style="list-style-type: none"> • There is a random data generator (1 pt.) • The map-reduce queries are correctly implemented (1 pts.) • The student makes use of design patterns to change the database layer and recycle the code for the front end application (1 pt.)



Assignment 3a	<ul style="list-style-type: none">• The graph database structure is correctly implemented (1 pt.)• The data respects the given constraints (1 pt.)• The queries in cypher are correctly implemented (1 pts.)
Assignment 3b	<ul style="list-style-type: none">• The graph database structure is correctly implemented (1 pt.)• The queries are correctly implemented (1 pt.)• Enough entities and relationships have been mapped into the database (1 pt.)



Exam structure

What follows is the general structure of a DEV5 exam. You can find a concrete exam sample in the course page on natschool.

Associated learning goals: RDBMS.

4.0.0.1 Question I: SQL queries

Given a database definition writes the SQL code for the requested queries

Associated learning goals: NORM, RDBMS.

4.0.0.2 Question II: Normalization

Given a relational schema and defined functional dependencies, provide a normalized version in BCNF

Associated learning goals: TRANS, RDBMS.

4.0.0.3 Question III: Transaction management and concurrency

Given the following N-queries, which are run in parallel, show a strict 2PL scheduling.

Associated learning goals: NONREL.

4.0.0.4 Question IV: Map-reduce

Implement the given queries with the map-reduce paradigm.

Associated learning goals: NONREL.

4.0.0.5 Question V: NoSQL databases en Graph Theory

Transform a given relational data model into a graph model and implement the given queries in Cypher